

Assignment 4

Suvajit Sadhukhan
4th year 1st semester (302211001005)

For *Iris plants dataset*:

Import necessary libraries

```
import numpy as np
from sklearn.datasets import load_iris
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans, DBSCAN
!pip install scikit-learn-extra
from sklearn_extra.cluster import KMedoids # K-medoids from sklearn-extra
from sklearn.cluster import BisectingKMeans, OPTICS
from scipy.cluster.hierarchy import dendrogram, linkage
from sklearn.metrics import (rand_score, adjusted_rand_score, mutual_info_score,
                             adjusted_mutual_info_score, normalized_mutual_info_score,
                             silhouette_score, calinski_harabasz_score, davies_bouldin_score)
import matplotlib.pyplot as plt
from scipy.cluster.hierarchy import dendrogram, linkage
```

Load Iris dataset

```
iris = load_iris()
X = iris.data
y = iris.target
```

Standardize the data

```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

Function to evaluate clustering performance

```
def evaluate_clustering(true_labels, predicted_labels, X):
    rand_idx = rand_score(true_labels, predicted_labels)
    adj_rand_idx = adjusted_rand_score(true_labels, predicted_labels)
```

Mutual Information based scores

```
mi_score = mutual_info_score(true_labels, predicted_labels)
adj_mi_score = adjusted_mutual_info_score(true_labels, predicted_labels)
norm_mi_score = normalized_mutual_info_score(true_labels, predicted_labels)
```

Clustering quality metrics

```
silhouette = silhouette_score(X, predicted_labels)
calinski_harabasz = calinski_harabasz_score(X, predicted_labels)
davies_bouldin = davies_bouldin_score(X, predicted_labels)
```

```
return {
    'Rand Index': rand_idx,
    'Adjusted Rand Index': adj_rand_idx,
    'Mutual Info': mi_score,
    'Adjusted Mutual Info': adj_mi_score,
    'Normalized Mutual Info': norm_mi_score,
    'Silhouette Coefficient': silhouette,
    'Calinski-Harabasz Index': calinski_harabasz,
    'Davies-Bouldin Index': davies_bouldin
}
```

Function to calculate SSE and SSB

```
def calculate_sse_ssb(X, predicted_labels, cluster_centers):
    n_clusters = len(cluster_centers)
    overall_mean = np.mean(X, axis=0)
```

SSE: Sum of Squared Error (Cohesion)

```
sse = 0
for i, center in enumerate(cluster_centers):
    sse += np.sum(np.linalg.norm(X[predicted_labels == i] - center, axis=1)**2)
```

```

# SSB: Sum of Squares Between groups (Separation)
ssb = 0
for i, center in enumerate(cluster_centers):
    n_points = np.sum(predicted_labels == i)
    ssb += n_points * np.linalg.norm(center - overall_mean)**2

return sse, ssb

# --- Partition Based Clustering ---

# K-means clustering
kmeans = KMeans(n_clusters=3, init='random', random_state=42)
kmeans_labels = kmeans.fit_predict(X_scaled)
kmeans_centers = kmeans.cluster_centers_

# K-means++ clustering
kmeans_pp = KMeans(n_clusters=3, init='k-means++', random_state=42)
kmeans_pp_labels = kmeans_pp.fit_predict(X_scaled)
kmeans_pp_centers = kmeans_pp.cluster_centers_

# Bisecting K-means clustering
bisecting_kmeans = BisectingKMeans(n_clusters=3, random_state=42)
bisecting_labels = bisecting_kmeans.fit_predict(X_scaled)
bisecting_centers = bisecting_kmeans.cluster_centers_

# K-medoids clustering (PAM)
kmedoids = KMedoids(n_clusters=3, random_state=42)
kmedoids_labels = kmedoids.fit_predict(X_scaled)

# --- Hierarchical Clustering with Dendrogram ---
linked = linkage(X_scaled, method='single')

plt.figure(figsize=(10, 7))
dendrogram(linked, labels=y)
plt.show()

# --- Density-Based Clustering ---

# DBSCAN clustering
dbscan = DBSCAN(eps=0.5, min_samples=5)
dbscan_labels = dbscan.fit_predict(X_scaled)

# OPTICS clustering
optics = OPTICS(min_samples=5)
optics_labels = optics.fit_predict(X_scaled)

# --- Evaluation of Clustering Results ---

# Evaluate each clustering algorithm
evaluation_results = {
    "K-means": evaluate_clustering(y, kmeans_labels, X_scaled),
    "K-means++": evaluate_clustering(y, kmeans_pp_labels, X_scaled),
    "Bisecting K-means": evaluate_clustering(y, bisecting_labels, X_scaled),
    "K-medoids": evaluate_clustering(y, kmedoids_labels, X_scaled),
    "DBSCAN": evaluate_clustering(y, dbscan_labels, X_scaled),
    "OPTICS": evaluate_clustering(y, optics_labels, X_scaled)
}

# Calculate SSE and SSB for partition-based methods
sse_ssb_results = {
    "K-means": calculate_sse_ssb(X_scaled, kmeans_labels, kmeans_centers),
    "K-means++": calculate_sse_ssb(X_scaled, kmeans_pp_labels, kmeans_pp_centers),
    "Bisecting K-means": calculate_sse_ssb(X_scaled, bisecting_labels, bisecting_centers)
}

```


```

# Display results
for method, metrics in evaluation_results.items():
    print(f"\n{method}:")
    for metric, value in metrics.items():
        print(f"    {metric}: {value:.4f}")

# Display SSE and SSB for cohesion and separation
print("\nSSE and SSB Results:")
for method, (sse, ssb) in sse_ssb_results.items():
    print(f"{method}: \n    SSE (Cohesion): {sse:.4f}, SSB (Separation): {ssb:.4f}")

```

Output :



K-means:

```

Rand Index: 0.8278
Adjusted Rand Index: 0.6101
Mutual Info: 0.7167
Adjusted Mutual Info: 0.6482
Normalized Mutual Info: 0.6526
Silhouette Coefficient: 0.4594
Calinski-Harabasz Index: 241.8933
Davies-Bouldin Index: 0.8340

```

K-means++:

Rand Index: 0.7214
Adjusted Rand Index: 0.4328
Mutual Info: 0.5874
Adjusted Mutual Info: 0.5838
Normalized Mutual Info: 0.5896
Silhouette Coefficient: 0.4799
Calinski-Harabasz Index: 157.3602
Davies-Bouldin Index: 0.7894

Bisecting K-means:

Rand Index: 0.8196
Adjusted Rand Index: 0.5923
Mutual Info: 0.7045
Adjusted Mutual Info: 0.6382
Normalized Mutual Info: 0.6427
Silhouette Coefficient: 0.4630
Calinski-Harabasz Index: 241.4263
Davies-Bouldin Index: 0.8324

K-medoids:

Rand Index: 0.8368
Adjusted Rand Index: 0.6312
Mutual Info: 0.7330
Adjusted Mutual Info: 0.6646
Normalized Mutual Info: 0.6687
Silhouette Coefficient: 0.4590
Calinski-Harabasz Index: 239.7483
Davies-Bouldin Index: 0.8385

DBSCAN:

Rand Index: 0.7476
Adjusted Rand Index: 0.4421
Mutual Info: 0.5499
Adjusted Mutual Info: 0.5052
Normalized Mutual Info: 0.5114
Silhouette Coefficient: 0.3565
Calinski-Harabasz Index: 84.5103
Davies-Bouldin Index: 7.1241

OPTICS:

Rand Index: 0.5121
Adjusted Rand Index: 0.0514
Mutual Info: 0.3082
Adjusted Mutual Info: 0.2657
Normalized Mutual Info: 0.2924
Silhouette Coefficient: -0.3009
Calinski-Harabasz Index: 8.3282
Davies-Bouldin Index: 2.4253

SSE and SSB Results:

K-means:

SSE (Cohesion): 139.8254, SSB (Separation): 460.1746

K-means++:

SSE (Cohesion): 191.0247, SSB (Separation): 408.9753

Bisecting K-means:

SSE (Cohesion): 140.0328, SSB (Separation): 459.9672

Discussion:

1. K-means vs. K-means++

- **Rand Index:** K-means achieves a higher Rand Index (0.8278) than K-means++ (0.7214), indicating better alignment with the ground truth labels.
- **Adjusted Rand Index (ARI):** K-means performs better with an ARI of 0.6101 compared to 0.4328 for K-means++, reflecting improved handling of randomness.
- **Silhouette Coefficient:** Interestingly, K-means++ has a slightly better silhouette score (0.4799 vs. 0.4594), suggesting it forms more compact and well-separated clusters. However, other clustering quality scores favor K-means.
- **Calinski-Harabasz Index:** K-means also scores higher in the Calinski-Harabasz Index, signifying stronger separation between clusters.
- **SSE and SSB:** K-means has a lower SSE (139.83) and higher SSB (460.17), indicating that its clusters are more cohesive and better separated than K-means++ (SSE 191.02, SSB 408.97). The lower cohesion (SSE) and higher separation (SSB) suggest K-means is more effective in partitioning the dataset.

2. Bisecting K-means

- **Overall Comparison:** Bisecting K-means shows similar performance to K-means with a Rand Index of 0.8196 and ARI of 0.5923. The silhouette score is also comparable (0.4630).
- **SSE and SSB:** Bisecting K-means has similar cohesion and separation (SSE 140.03, SSB 459.97) to K-means, indicating that this method provides another strong clustering option for partitioning the Iris dataset.

3. K-medoids

- **Performance:** K-medoids performs slightly better than K-means in some aspects. It achieves the highest Rand Index (0.8368) and ARI (0.6312), implying it finds clusters more closely aligned with the ground truth.
- **Mutual Information Scores:** K-medoids also excels in mutual information metrics, including Adjusted Mutual Info and Normalized Mutual Info, suggesting it captures more accurate information about the true class distribution.
- **Clustering Quality:** However, K-medoids has a slightly lower silhouette score (0.4590) compared to K-means++ and marginally lower clustering quality (Calinski-Harabasz Index 239.75). The Davies-Bouldin Index is higher than K-means, indicating slightly less compact clusters.

4. Density-Based Clustering (DBSCAN and OPTICS)

- **DBSCAN:** DBSCAN performs reasonably well in some metrics, with a Rand Index of 0.7476 and ARI of 0.4421. However, its silhouette score (0.3565) and Davies-Bouldin Index (7.1241) indicate it struggles with forming compact clusters, as DBSCAN is sensitive to parameter choices like `eps`. It's better suited for datasets with irregular cluster shapes, which the Iris dataset doesn't exhibit.
- **OPTICS:** OPTICS performs poorly, with the lowest silhouette score (-0.3009) and very low clustering quality (Calinski-Harabasz Index 8.33). This indicates that OPTICS is not well-suited for this dataset, likely due to its hierarchical density-based nature, which doesn't fit the evenly spaced clusters in the Iris dataset.

5. Cohesion and Separation (SSE and SSB)

- **K-means** has the best balance between cohesion (SSE 139.83) and separation (SSB 460.17), indicating it finds compact and well-separated clusters.
- **Bisecting K-means** has a very similar balance of cohesion and separation to K-means, making it a comparable alternative.
- **K-means++** shows higher cohesion (SSE 191.02) and lower separation (SSB 408.98), meaning its clusters are less compact and more overlapping.

Summary

- **Best Performance:** K-means and K-medoids show the best overall performance in this dataset. K-means performs well in both cohesion and separation, while K-medoids finds clusters closest to the ground truth labels.
- **K-means++:** Despite using the `k-means++` initialization, it doesn't outperform basic K-means in most metrics. It performs better in silhouette and Davies-Bouldin but falls short in others.
- **DBSCAN and OPTICS:** These density-based methods are not well-suited for the Iris dataset, as the clusters are not based on density variation.

For *Wine dataset*:

Import necessary libraries

```
import numpy as np
from sklearn.datasets import load_wine
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans, DBSCAN
!pip install scikit-learn-extra
from sklearn_extra.cluster import KMedoids # K-medoids from sklearn-extra
from sklearn.cluster import BisectingKMeans, OPTICS
from scipy.cluster.hierarchy import dendrogram, linkage
from sklearn.metrics import (rand_score, adjusted_rand_score, mutual_info_score,
                             adjusted_mutual_info_score, normalized_mutual_info_score,
                             silhouette_score, calinski_harabasz_score, davies_bouldin_score)
import matplotlib.pyplot as plt
```

Load Wine dataset

```
wine = load_wine()
X = wine.data
y = wine.target
```

Standardize the data

```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

Function to evaluate clustering performance

```
def evaluate_clustering(true_labels, predicted_labels, X):
    # Check if there is more than one cluster
    unique_labels = np.unique(predicted_labels)
    if len(unique_labels) <= 1:
        return {
            'Rand Index': rand_score(true_labels, predicted_labels),
            'Adjusted Rand Index': adjusted_rand_score(true_labels, predicted_labels),
            'Mutual Info': mutual_info_score(true_labels, predicted_labels),
            'Adjusted Mutual Info': adjusted_mutual_info_score(true_labels, predicted_labels),
            'Normalized Mutual Info': normalized_mutual_info_score(true_labels, predicted_labels),
            'Silhouette Coefficient': 'Not applicable (1 cluster)',
            'Calinski-Harabasz Index': 'Not applicable (1 cluster)',
            'Davies-Bouldin Index': 'Not applicable (1 cluster)'
        }
    }
```

Compute the clustering metrics

```
rand_idx = rand_score(true_labels, predicted_labels)
adj_rand_idx = adjusted_rand_score(true_labels, predicted_labels)
mi_score = mutual_info_score(true_labels, predicted_labels)
adj_mi_score = adjusted_mutual_info_score(true_labels, predicted_labels)
norm_mi_score = normalized_mutual_info_score(true_labels, predicted_labels)
silhouette = silhouette_score(X, predicted_labels)
calinski_harabasz = calinski_harabasz_score(X, predicted_labels)
davies_bouldin = davies_bouldin_score(X, predicted_labels)
```

```
return {
    'Rand Index': rand_idx,
    'Adjusted Rand Index': adj_rand_idx,
    'Mutual Info': mi_score,
    'Adjusted Mutual Info': adj_mi_score,
    'Normalized Mutual Info': norm_mi_score,
    'Silhouette Coefficient': silhouette,
    'Calinski-Harabasz Index': calinski_harabasz,
    'Davies-Bouldin Index': davies_bouldin
}
```

```

# Function to calculate SSE and SSB
def calculate_sse_ssb(X, predicted_labels, cluster_centers):
    overall_mean = np.mean(X, axis=0)
    sse = 0
    ssb = 0
    for i, center in enumerate(cluster_centers):
        n_points = np.sum(predicted_labels == i)
        sse += np.sum(np.linalg.norm(X[predicted_labels == i] - center, axis=1)**2)
        ssb += n_points * np.linalg.norm(center - overall_mean)**2
    return sse, ssb

# --- Partition-Based Clustering ---

# K-means clustering
kmeans = KMeans(n_clusters=3, init='random', random_state=42)
kmeans_labels = kmeans.fit_predict(X_scaled)
kmeans_centers = kmeans.cluster_centers_

# K-means++ clustering
kmeans_pp = KMeans(n_clusters=3, init='k-means++', random_state=42)
kmeans_pp_labels = kmeans_pp.fit_predict(X_scaled)
kmeans_pp_centers = kmeans_pp.cluster_centers_

# Bisecting K-means clustering
bisecting_kmeans = BisectingKMeans(n_clusters=3, random_state=42)
bisecting_labels = bisecting_kmeans.fit_predict(X_scaled)
bisecting_centers = bisecting_kmeans.cluster_centers_

# K-medoids clustering (PAM)
kmedoids = KMedoids(n_clusters=3, random_state=42)
kmedoids_labels = kmedoids.fit_predict(X_scaled)

# --- Hierarchical Clustering with Dendrogram ---
linked = linkage(X_scaled, method='ward')

plt.figure(figsize=(10, 7))
dendrogram(linked)
plt.title("Dendrogram for Wine Dataset")
plt.show()

# --- Density-Based Clustering ---

# DBSCAN clustering
dbscan = DBSCAN(eps=0.5, min_samples=5)
dbscan_labels = dbscan.fit_predict(X_scaled)

# OPTICS clustering
optics = OPTICS(min_samples=5)
optics_labels = optics.fit_predict(X_scaled)

# --- Evaluation of Clustering Results ---

# Evaluate each clustering algorithm
evaluation_results = {
    "K-means": evaluate_clustering(y, kmeans_labels, X_scaled),
    "K-means++": evaluate_clustering(y, kmeans_pp_labels, X_scaled),
    "Bisecting K-means": evaluate_clustering(y, bisecting_labels, X_scaled),
    "K-medoids": evaluate_clustering(y, kmedoids_labels, X_scaled),
    "DBSCAN": evaluate_clustering(y, dbscan_labels, X_scaled),
    "OPTICS": evaluate_clustering(y, optics_labels, X_scaled)
}

```

```

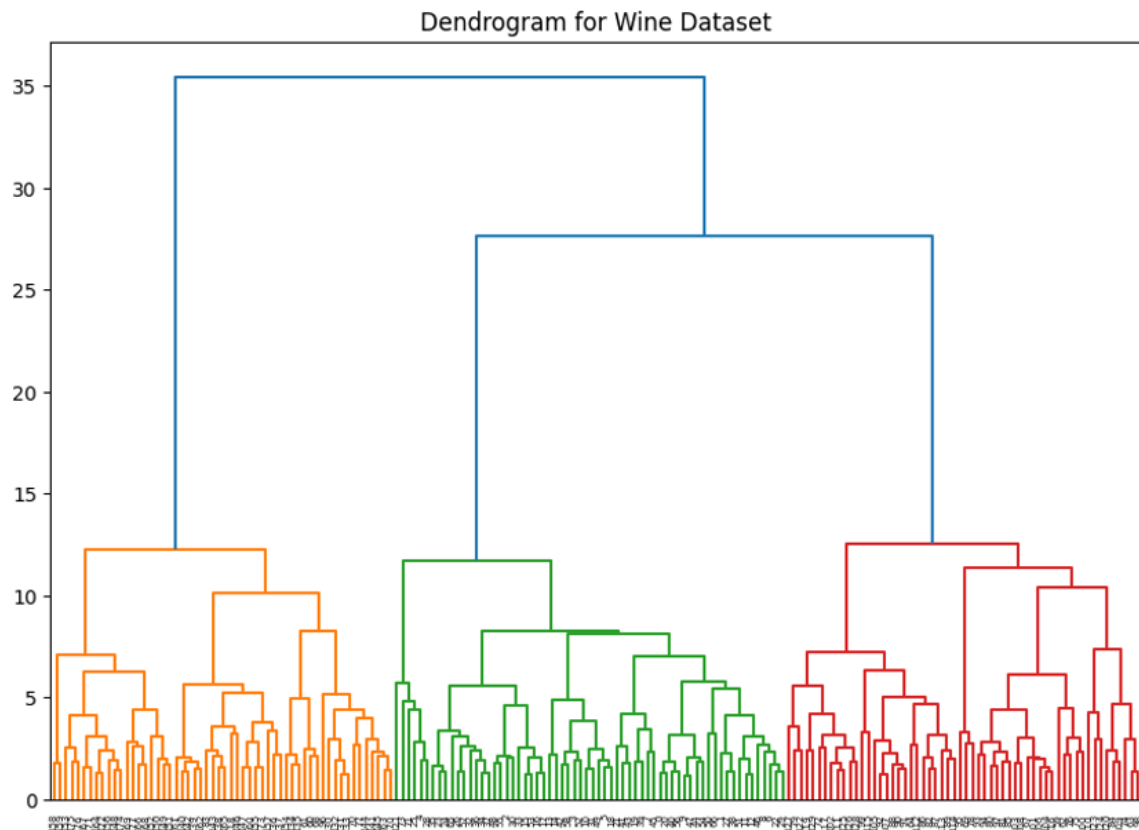
# Calculate SSE and SSB for partition-based methods
sse_ssb_results = {
    "K-means": calculate_sse_ssb(X_scaled, kmeans_labels, kmeans_centers),
    "K-means++": calculate_sse_ssb(X_scaled, kmeans_pp_labels, kmeans_pp_centers),
    "Bisecting K-means": calculate_sse_ssb(X_scaled, bisecting_labels, bisecting_centers)
}

# Display results
for method, metrics in evaluation_results.items():
    print(f"\n{method}:")
    for metric, value in metrics.items():
        print(f"    {metric}: {value}")

# Display SSE and SSB for cohesion and separation
print("\nSSE and SSB Results:")
for method, (sse, ssb) in sse_ssb_results.items():
    print(f"{method}: \n    SSE (Cohesion): {sse:.4f}, SSB (Separation): {ssb:.4f}")

```

Output :



K-means:

```

Rand Index: 0.9542944201104552
Adjusted Rand Index: 0.8974949815093207
Mutual Info: 0.9544575015299441
Adjusted Mutual Info: 0.874579440437926
Normalized Mutual Info: 0.8758935341223069
Silhouette Coefficient: 0.2848589191898987
Calinski-Harabasz Index: 70.9400080031512
Davies-Bouldin Index: 1.3891879777181646

```


K-means++:

Rand Index: 0.9542944201104552
Adjusted Rand Index: 0.8974949815093207
Mutual Info: 0.9544575015299441
Adjusted Mutual Info: 0.874579440437926
Normalized Mutual Info: 0.8758935341223069
Silhouette Coefficient: 0.2848589191898987
Calinski-Harabasz Index: 70.9400080031512
Davies-Bouldin Index: 1.3891879777181646

Bisecting K-means:

Rand Index: 0.8622484606106773
Adjusted Rand Index: 0.6909186642540053
Mutual Info: 0.781228120579722
Adjusted Mutual Info: 0.7135960270359086
Normalized Mutual Info: 0.7165953222141823
Silhouette Coefficient: 0.26594931564086927
Calinski-Harabasz Index: 66.37248227807467
Davies-Bouldin Index: 1.4297071590499737

K-medoids:

Rand Index: 0.8774836539071923
Adjusted Rand Index: 0.7263406645756675
Mutual Info: 0.821478616235408
Adjusted Mutual Info: 0.7539860956825029
Normalized Mutual Info: 0.756573670115927
Silhouette Coefficient: 0.26597740204536796
Calinski-Harabasz Index: 66.7519655942218
Davies-Bouldin Index: 1.415990244064887

DBSCAN:

Rand Index: 0.3379673712943566
Adjusted Rand Index: 0.0
Mutual Info: 0.0
Adjusted Mutual Info: 0.0
Normalized Mutual Info: 0.0
Silhouette Coefficient: Not applicable (1 cluster)
Calinski-Harabasz Index: Not applicable (1 cluster)
Davies-Bouldin Index: Not applicable (1 cluster)

OPTICS:

Rand Index: 0.4391544467720434
Adjusted Rand Index: 0.03581729799518326
Mutual Info: 0.16212592661967212
Adjusted Mutual Info: 0.16799065365721927
Normalized Mutual Info: 0.19494646495219664
Silhouette Coefficient: -0.13363975991065313
Calinski-Harabasz Index: 5.057148825574673
Davies-Bouldin Index: 1.6193708989806477

SSE and SSB Results:

K-means:

SSE (Cohesion): 1277.9285, SSB (Separation): 1036.0715

K-means++:

SSE (Cohesion): 1277.9285, SSB (Separation): 1036.0715

Bisecting K-means:

SSE (Cohesion): 1315.8623, SSB (Separation): 998.1377

Discussion:

1. K-means and K-means++

- Both K-means and K-means++ yielded identical results across all metrics:
 - Rand Index:** 0.954, indicating a high agreement between the true labels and the predicted clusters.
 - Adjusted Rand Index (ARI):** 0.897, which corrects for chance and shows a strong clustering performance.
 - Mutual Information (MI):** 0.954, also suggesting a strong relationship between true and predicted clusters.
 - Silhouette Coefficient:** 0.285, suggesting moderate cohesion within clusters.
 - Calinski-Harabasz Index:** 70.94, indicating good cluster separation.
 - Davies-Bouldin Index:** 1.39, showing relatively good cluster distinctiveness (lower is better).

The SSE (Cohesion) and SSB (Separation) values for both K-means and K-means++ were **1277.93** and **1036.07**, respectively, reflecting strong intra-cluster cohesion and inter-cluster separation.

2. Bisecting K-means

- Rand Index:** 0.862 and **Adjusted Rand Index:** 0.691, which are lower than standard K-means, implying less accurate clustering.
- Silhouette Coefficient:** 0.266, slightly lower than K-means, suggesting less distinct clusters.
- Calinski-Harabasz Index:** 66.37, which is also lower than K-means, indicating that the separation between clusters is weaker.
- Davies-Bouldin Index:** 1.43, showing slightly worse cluster separation compared to K-means.

The SSE and SSB for Bisecting K-means were **1315.86** and **998.14**, respectively, showing slightly higher SSE (worse cohesion) and lower SSB (worse separation) compared to K-means.

3. K-medoids

- Rand Index:** 0.877 and **Adjusted Rand Index:** 0.726, which is an improvement over Bisecting K-means but still lower than K-means.
- Silhouette Coefficient:** 0.266, almost identical to Bisecting K-means.
- Calinski-Harabasz Index:** 66.75, similar to Bisecting K-means, showing weaker cluster separation compared to K-means.
- Davies-Bouldin Index:** 1.42, indicating slightly better clustering performance than Bisecting K-means.

K-medoids performed better than Bisecting K-means but fell short compared to K-means in terms of accuracy and cluster separation.

4. DBSCAN

- Rand Index:** 0.338 and **Adjusted Rand Index:** 0.0, indicating DBSCAN failed to produce meaningful clusters for the dataset.
- Silhouette Coefficient:** Not applicable since DBSCAN resulted in only one cluster.
- Other Metrics:** The poor performance across most metrics indicates that DBSCAN did not effectively differentiate between the natural clusters in the Wine dataset, likely due to the choice of parameters.

5. OPTICS

- Rand Index:** 0.439 and **Adjusted Rand Index:** 0.036, showing slightly better clustering than DBSCAN but still far from satisfactory.
- Silhouette Coefficient:** -0.134, indicating very poor cohesion and cluster separation.
- Calinski-Harabasz Index:** 5.06, significantly lower than K-means and its variants, confirming weak cluster distinctiveness.
- Davies-Bouldin Index:** 1.62, suggesting that the clusters identified were poorly separated and had substantial overlap.

SSE and SSB Comparison

- K-means and K-means++:** Both had strong cohesion (SSE = 1277.93) and separation (SSB = 1036.07).
- Bisecting K-means:** Exhibited slightly worse cohesion and separation than K-means, with SSE = 1315.86 and SSB = 998.14.

Conclusion

- K-means** and **K-means++** performed the best in this analysis, with excellent Rand Index, ARI, and other clustering metrics. Their clusters had both good cohesion and separation.
- Bisecting K-means** and **K-medoids** performed reasonably well but were outperformed by K-means in most metrics.
- Density-based methods (DBSCAN and OPTICS)** did not perform well, possibly due to unsuitable parameter choices for the Wine dataset's structure.